

Cliente Exemplo S.A.

Score Justification

Detailed Methodology and Calculation Rationale

ASSESSMENT	AI Maturity Assessment 2026 H1
INDUSTRY	Financial Services / Insurance
ASSESSMENT DATE	2026-05-08
FRAMEWORK VERSION	1.0.0
GENERATED ON	2026-05-08
SCOPE	Engineering Organization

Table of Contents

1. Executive Overview	.
1.1 Purpose of This Document	.
1.2 Summary of Final Scores	.
2. Scoring Methodology	.
2.1 Maturity Level Definitions	.
2.2 Scoring Guidelines	.
2.3 Weighted Average Calculation Formula	.
3. P1: Developer Productivity — Detailed Scores	.
4. P2: DevOps Lifecycle — Detailed Scores	.
5. P3: Application Platform — Detailed Scores	.
6. Overall Maturity Calculation	.
7. Key Evidence Sources	.
8. Conclusion	.

1. Executive Overview

This document provides a comprehensive explanation of the scoring methodology used in the Cliente Exemplo S.A. AI Maturity Assessment, including detailed justifications for each capability score and the weighted average calculations that determine pillar and overall maturity levels.

1.1 Purpose of This Document

- Provide full transparency on how each score was determined.
- Document the evidence base used for scoring decisions.
- Explain the weighted average calculation methodology.
- Justify the rationale behind each maturity level assignment.
- Enable stakeholders to understand and validate assessment results.

1.2 Summary of Final Scores

PILLAR	WEIGHTED AVG	MATURITY LEVEL	TARGET	GAP
P1: Developer Productivity	2.69	L3 — Gerenciado	L4.0	1.3
P2: DevOps Lifecycle	1.52	L2 — Definido	L4.0	2.5
P3: Application Platform	1.92	L2 — Definido	L4.0	2.1
OVERALL	1.99	L2 — Definido	L4.0	2.0



2. Scoring Methodology

2.1 Maturity Level Definitions

The AI Maturity Framework uses a 5-level scale (L0-L4) to measure capability maturity. Each level has specific characteristics and evidence requirements.

LEVEL	NAME	DESCRIPTION	EVIDENCE REQUIRED
L0	Traditional	No structured approach. Manual processes dominate. No AI tooling adoption.	Observation of manual workflows, absence of automation tools
L1	AI-Assisted	Beginning to adopt AI tools. Pilots underway (<20% adoption). Basic automation exists.	License counts, pilot team lists, initial metrics
L2	AI-Enhanced	Significant AI adoption (50%+). Comprehensive automation. Self-service emerging.	Adoption dashboards, automated workflow logs, self-service portal
L3	AI-Optimized	Wide AI adoption (75%+). AI agents for complex tasks. Predictive capabilities.	AI usage reports, agent task logs, predictive analytics dashboards
L4	Agentic / Autonomous	Full AI integration. Multi-agent systems. Self-healing with minimal human intervention.	Autonomous operation logs, multi-agent coordination evidence, MTTR metrics

2.2 Scoring Guidelines

The following principles were applied consistently during scoring:


- Score based on EVIDENCE, not aspirations — always require proof.
- When capabilities span multiple levels, score at the LOWER level.
- Consider CONSISTENCY across teams — isolated excellence does not raise the org score.
- Platform Engineering [PE] capabilities indicate Agentic Platform Engineering readiness.
- Use half-points (e.g. 1.5) only when strong evidence supports partial progress.

2.3 Weighted Average Calculation Formula

Each capability has an assigned weight (2 or 3) based on its strategic importance. The Pillar score is calculated as:

Pillar Score = $\Sigma(\text{capability_score} \times \text{weight}) / \Sigma(\text{weights})$
Overall Score = average(pillar_scores)

3. P1: Developer Productivity — Detailed Scores



3.1 Pillar Summary

Total Capabilities	8
Sum of Weights	18.00
Weighted Sum of Scores	22 (3×1.0 + 2×1.0 + 2×1.5 + 2×1.5 + 2×1.0 + 3×1.0 + 2×1.0 + 2×2.0)
Weighted Average	2.69 (= weighted sum / 18.00)
Maturity Level	L3 — Gerenciado

3.2 Capability-by-Capability Justification

PE 1.1 AI Coding Assistants & Adoption

SCORE	WEIGHT	PE INDICATOR
L3 — Gerenciado	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- 645 AI coding licenses deployed (out of 2,430 developers)
- Penetration rate: 26.5%
- 30% suggestion acceptance rate
- 0 Agent Mode usage observed

SCORING RATIONALE

Score of L1 was assigned because the organization is piloting AI coding assistants with approximately 26% of developers, well below the 50% threshold required for L2. While initial deployment exists, suggestion acceptance is below target and advanced features (Agent Mode, agentic workflows) show zero usage.

1.2 Code Review with AI developer group

SCORE	WEIGHT	PE INDICATOR
L2 — Definido	2 (Medium)	No

EVIDENCE COLLECTED

- AI code review available but not standardized
- <10% of pull requests use AI-assisted review
- No policy mandates AI review usage
- Manual review remains primary mechanism

SCORING RATIONALE

Score of L1 reflects nascent AI code review usage. Tools are available to all developers but adoption is sporadic, with fewer than 10% of PRs leveraging AI assistance. Without a standardized policy or measurable enforcement, the capability remains in pilot phase.

1.3 Test Generation & TDD with AI developer group

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	2 (Medium)	No

EVIDENCE COLLECTED

- AI test generation in pilot for 2 teams
- Code coverage stable at 80%
- No mutation testing in place
- Test maintenance still primarily manual

SCORING RATIONALE

Score of L1.5 reflects piloted but inconsistent AI test generation. Two teams have evidence of consistent usage with measurable coverage improvements, but the practice is not organization-wide. The half-point above L1 acknowledges concrete pilot success.

1.4 Documentation Automation developer group

SCORE	WEIGHT	PE INDICATOR
L3 — Gerenciado	2 (Medium)	No

EVIDENCE COLLECTED

- Some teams use AI for inline documentation
- API documentation auto-generated for 30% of services
- Architecture decision records still manual
- No standardized doc-generation policy

SCORING RATIONALE

Score of L1.5 reflects partial automation: API docs auto-generated for a subset of services, but ADRs and runbooks still manual. The half-point recognizes meaningful automation in production for code-level documentation.

1.5 Developer Onboarding & Skills

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	2 (Medium)	No

EVIDENCE COLLECTED

- Developer onboarding takes 4-6 weeks to first commit
- Mentorship is informal
- AI tools introduced via 1-hour orientation
- No measurable AI skills assessment

SCORING RATIONALE

Score of L1 because onboarding includes basic AI tool exposure but lacks structured ramp-up paths, formal AI skills tracks, or measurable proficiency gates. The 4-6 week time-to-first-commit is well above industry benchmarks.

PE 1.6 IDE & Local Dev Environment developer group

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- IDE setup takes 4-8 hours per developer
- No standardized dev container configuration
- Cloud IDE in pilot for 1 team
- Local environment drift causes 20% of new-dev support tickets

SCORING RATIONALE

Score of L1 reflects the manual, drift-prone state of local environments. While a cloud IDE pilot exists, it is not the default. Setup time and environment-related support tickets indicate significant opportunity for standardization.

1.7 AI-Assisted Refactoring & Modernization

developer group

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	2 (Medium)	No

EVIDENCE COLLECTED

- Manual refactoring is the norm
- Legacy modernization done case-by-case
- No AI-driven dependency upgrade automation
- Technical debt tracking exists but not AI-prioritized

SCORING RATIONALE

Score of L1 reflects awareness of refactoring needs but lack of AI tooling integration. The organization tracks technical debt manually but has no automated prioritization or AI-assisted modernization in production.

1.8 Productivity Metrics & SPACE

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	2 (Medium)	No

EVIDENCE COLLECTED

- DORA metrics tracked at portfolio level
- SPACE framework partially adopted (3 of 5 dimensions)
- Quarterly developer experience surveys conducted
- No real-time productivity dashboard

SCORING RATIONALE

Score of L2 reflects established DORA tracking and partial SPACE adoption. Surveys provide qualitative signal, but real-time AI-driven productivity insights are absent.

3.3 Calculation Summary

Pillar Score = 22 (3×1.0 + 2×1.0 + 2×1.5 + 2×1.5 + 2×1.0 + 3×1.0 + 2×1.0 + 2×2.0) / 18.00

= 2.69 → L3 – Gerenciado

4. P2: DevOps Lifecycle – Detailed Scores

P2: DevOps Lifecycle

L1.5 → **L4.0**

Gap +2.48



4.1 Pillar Summary

Total Capabilities	10
Sum of Weights	26.00
Weighted Sum of Scores	47 (3×1.5 + 3×2.0 + 3×2.0 + 2×2.0 + 3×2.5 + 2×2.0 + 3×2.5 + 2×1.5 + 2×0.0 + 3×2.0)
Weighted Average	1.52 (= weighted sum / 26.00)
Maturity Level	L2 — Definido

4.2 Capability-by-Capability Justification

2.1 Source Control & Branching Strategy

SCORE	WEIGHT	PE INDICATOR
L2 — Definido	3 (High - Critical)	No

EVIDENCE COLLECTED

- Git-based source control across all teams
- Branching strategies vary (GitFlow vs trunk-based)
- No org-wide branch policy automation
- Some repos enforce signed commits

SCORING RATIONALE

Score of L1.5 reflects modern source control with inconsistent branching practices. The half-point above L1 recognizes signed-commit enforcement on a meaningful subset of repositories.

PE 2.2 CI/CD Pipelines

SCORE	WEIGHT	PE INDICATOR
Sem respuesta	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- All teams have automated CI
- CD pipelines exist for 60% of services
- Pipeline-as-code in widespread use
- Manual gates remain in production deployments

SCORING RATIONALE

Score of L2 reflects mature CI with partial CD coverage. Pipeline-as-code is standard, but the 40% of services without automated CD and the prevalence of manual production gates prevent advancement to L3.

2.3 Build & Artifact Management

SCORE	WEIGHT	PE INDICATOR
L2 — Definido	3 (High - Critical)	No

EVIDENCE COLLECTED

- Centralized artifact registry in use
- Build caching enabled for major repos
- SBOM generation for 50% of artifacts
- Reproducible builds not enforced org-wide

SCORING RATIONALE

Score of L2 reflects centralized artifact management with build optimization. SBOM coverage at 50% and lack of org-wide reproducible-build enforcement keep the score below L3.

2.4 Test Automation

SCORE	WEIGHT	PE INDICATOR
L1 — Em Desenvolvimento	2 (Medium)	No

EVIDENCE COLLECTED

- Unit test automation in CI for all services
- Integration tests for 70% of services
- E2E tests in nightly pipelines
- Flaky-test detection still manual

SCORING RATIONALE

Score of L2 reflects strong unit and integration test automation. The lack of automated flaky-test detection and incomplete E2E coverage in main pipelines prevent L3.

PE 2.5 Security & DevSecOps

SCORE	WEIGHT	PE INDICATOR
Sem resposta	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- SAST and SCA in 100% of pipelines
- DAST for 60% of public-facing services
- Secret scanning enforced
- Threat modeling tooling in pilot

SCORING RATIONALE

Score of L2.5 reflects mature shift-left security with broad SAST/SCA adoption and enforced secret scanning. The half-point above L2 acknowledges DAST for the majority of public services and threat-modeling pilots.

2.6 Release & Deployment Strategy

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- Blue-green deployments for critical services
- Canary releases in 30% of services
- Feature flags for 40% of new features
- Automated rollback for 50% of deployments

SCORING RATIONALE

Score of L2 reflects modern deployment strategies in production but with uneven adoption. Canary at 30% and feature flags at 40% indicate progressive delivery is the exception rather than the rule.

PE 2.7 Observability & Monitoring

SCORE	WEIGHT	PE INDICATOR
Sem resposta	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- Distributed tracing for 80% of services
- Centralized logging org-wide
- SLOs defined for 70% of services
- Alert tuning still manual

SCORING RATIONALE

Score of L2.5 reflects strong observability foundations with distributed tracing, centralized logs, and SLO discipline. The half-point above L2 reflects breadth of tracing and SLO adoption. Manual alert tuning prevents L3.

2.8 Incident Management & SRE

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- On-call rotations exist for tier-1 services
- Postmortems written but not always actionable
- MTTR varies widely (30 min to 8 hours)
- No formal SRE practice

SCORING RATIONALE

Score of L1.5 reflects basic incident response with significant variability in MTTR. The half-point recognizes consistent on-call coverage for tier-1 services. Lack of formal SRE practice limits L2.

2.9 Chaos Engineering

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- No chaos engineering practice in place
- No game days conducted
- Disaster recovery exercises annual only
- No fault injection tooling deployed

SCORING RATIONALE

Score of L0 reflects complete absence of chaos engineering practices. While disaster recovery exercises occur annually, they do not constitute proactive chaos testing.

2.10 Compliance & Audit Automation

SCORE	WEIGHT	PE INDICATOR
L1 — Em Desenvolvimento	3 (High - Critical)	No

EVIDENCE COLLECTED

- Policy-as-code for cloud configurations
- Audit trails captured for 80% of changes
- Compliance reporting partially automated
- Manual evidence collection for major audits

SCORING RATIONALE

Score of L2 reflects established policy-as-code and broad audit-trail coverage. The manual nature of major audit evidence collection prevents L3.

4.3 Calculation Summary


Pillar Score = $47 (3 \times 1.5 + 3 \times 2.0 + 3 \times 2.0 + 2 \times 2.0 + 3 \times 2.5 + 2 \times 2.0 + 3 \times 2.5 + 2 \times 1.5 + 2 \times 0.0 + 3 \times 2.0) / 26.00$
= 1.52 → L2 — Definido

5. P3: Application Platform – Detailed Scores

P3: Application Platform

L1.9 → L4.0

Gap +2.08



5.1 Pillar Summary

Total Capabilities	10
Sum of Weights	26.00
Weighted Sum of Scores	55.5 (2×2.0 + 3×0.5 + 3×3.0 + 2×2.0 + 3×2.0 + 2×2.5 + 3×3.0 + 2×2.0 + 2×1.5 + 3×3.0)
Weighted Average	1.92 (= weighted sum / 26.00)
Maturity Level	L2 — Definido

5.2 Capability-by-Capability Justification

PE 3.1 Cloud-Native Architecture

SCORE	WEIGHT	PE INDICATOR
L3 — Gerenciado	2 (Medium)	Yes
EVIDENCE COLLECTED 70% of services containerized - Microservices for new builds - Legacy monoliths still in production for 30% of workloads - Service mesh in pilot		
SCORING RATIONALE Score of L2 reflects significant cloud-native adoption. The remaining 30% of monolithic workloads and the pilot stage of service mesh keep the score below L3.		

PE 3.2 Internal Developer Platform (IDP)

SCORE	WEIGHT	PE INDICATOR
Sem resposta	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- No formal IDP in place
- Self-service primarily through documented runbooks
- Service catalog in early planning phase
- 1 team experimenting with portal tooling

SCORING RATIONALE

Score of L0.5 reflects early-stage IDP thinking with no production deployment. The half-point recognizes the experimental work by one team. The complete absence of a service catalog or self-service portal keeps the score below L1.

3.3 Containerization & Orchestration

SCORE	WEIGHT	PE INDICATOR
L2 — Definido	3 (High - Critical)	No

EVIDENCE COLLECTED

- Kubernetes is the production runtime
- Multi-cluster federation in place
- Auto-scaling configured for 90% of services
- GitOps for cluster configuration

SCORING RATIONALE

Score of L3 reflects mature Kubernetes practice with multi-cluster federation, broad auto-scaling, and GitOps for configuration. The remaining gap to L4 is largely about AI-driven workload placement and predictive scaling.

3.4 API & Microservices Management

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- API gateway in production for external APIs
- OpenAPI specs for 70% of services
- API versioning conventions established
- API contract testing not enforced

SCORING RATIONALE

Score of L2 reflects mature API gateway with broad OpenAPI adoption and versioning discipline. The lack of enforced contract testing prevents L3.

3.5 Data Platform & Lakehouse

SCORE	WEIGHT	PE INDICATOR
L1 — Em Desenvolvimento	3 (High - Critical)	No

EVIDENCE COLLECTED

- Lakehouse architecture deployed
- ETL pipelines for major sources
- Data quality monitoring partial
- Data mesh principles in pilot for 1 domain

SCORING RATIONALE

Score of L2 reflects modern lakehouse architecture with established ETL. Partial data quality monitoring and the pilot stage of data mesh keep the score below L3.

PE 3.6 AI/ML Platform

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	Yes

EVIDENCE COLLECTED

- MLOps platform deployed for 5 use cases
- Model registry and versioning in place
- Feature store in pilot
- A/B testing for model deployments

SCORING RATIONALE

Score of L2.5 reflects mature MLOps for production use cases with versioning and registry. The half-point above L2 recognizes A/B testing for model deployments. Feature store pilot stage prevents L3.

3.7 Infrastructure-as-Code

SCORE	WEIGHT	PE INDICATOR
Sem resposta	3 (High - Critical)	No

EVIDENCE COLLECTED

- All infrastructure provisioned via IaC
- Terraform modules standardized
- Drift detection automated
- Module testing comprehensive

SCORING RATIONALE

Score of L3 reflects mature IaC with module standardization, drift detection, and comprehensive testing. The remaining gap to L4 concerns AI-assisted module generation and policy intelligence.

3.8 FinOps & Cost Optimization

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- Cost allocation tags enforced
- Monthly cost reviews per team
- Reserved capacity planning quarterly
- No automated rightsizing

SCORING RATIONALE

Score of L2 reflects mature cost discipline with allocation tags and review cadence. The absence of automated rightsizing and predictive cost analytics prevents L3.

3.9 Reliability & Resilience Engineering

SCORE	WEIGHT	PE INDICATOR
Sem resposta	2 (Medium)	No

EVIDENCE COLLECTED

- Multi-AZ deployments for tier-1 services
- Multi-region for 20% of services
- Backup strategies documented
- Recovery testing inconsistent

SCORING RATIONALE

Score of L1.5 reflects multi-AZ resilience with limited multi-region coverage. The half-point above L1 recognizes consistent multi-AZ for critical services. Inconsistent recovery testing prevents L2.

PE 3.10 Platform Adoption & Engineering Culture

SCORE	WEIGHT	PE INDICATOR
L3	3 (High - Critical)	Yes

EVIDENCE COLLECTED

- Quarterly engineering culture surveys
- Strong DevOps community of practice
- Internal tech talks and demo days
- Platform team has clear charter and product mindset

SCORING RATIONALE

Score of L3 reflects mature platform-as-product mindset with active community and feedback loops. The remaining gap to L4 concerns AI-driven platform evolution and predictive platform improvements.

5.3 Calculation Summary

Pillar Score = $55.5 (2 \times 2.0 + 3 \times 0.5 + 3 \times 3.0 + 2 \times 2.0 + 3 \times 2.0 + 2 \times 2.5 + 3 \times 3.0 + 2 \times 2.0 + 2 \times 1.5 + 3 \times 3.0) / 26.00$
= 1.92 → L2 — Definido

6. Overall Maturity Calculation

6.1 Pillar Score Summary

PILLAR	WEIGHTED AVG	MATURITY LEVEL
P1: Developer Productivity	2.69	L3 — Gerenciado
P2: DevOps Lifecycle	1.52	L2 — Definido
P3: Application Platform	1.92	L2 — Definido

6.2 Overall Score Calculation

Overall Score = $(2.69 + 1.52 + 1.92) / 3$
= 1.99 → L2 — Definido

6.3 Agentic Platform Engineering Readiness

2.60

PE Weighted Score

L3 — Gerenciado

AGENTIC PLATFORM ENGINEERING

Recommended Path: **Build foundation first**

PE Readiness score 1.48 (LOW) indicates basic capabilities exist but foundational gaps must close before adopting either Three Horizons (Red Hat Developer Hub) or Open Horizons (Backstage). A small Open Horizons pilot may be valuable to learn before scaling, but full IDP adoption should wait until score reaches MEDIUM (≥ 2.0).

SCORE RANGE	PE LEVEL	RECOMMENDED PATH	INTERPRETATION
< 1.0	NOT READY	Build foundation first	Significant foundational work needed before adopting either IDP path. Focus on standardization, source-of-truth, basic CI/CD.
1.0 – 1.9	LOW	Build foundation first	Basic capabilities exist. Close fundamental gaps before scaling. A small Open Horizons (Backstage) pilot may be valuable to learn before committing.
2.0 – 2.9	MEDIUM	Open Horizons (Backstage) / Three Horizons (Red Hat Developer Hub)	Foundation is in place. Open Horizons is the lower-risk default; Three Horizons fits organizations already in the Red Hat / OpenShift ecosystem.
≥ 3.0	HIGH	Either path — strategic choice	Either path can be chosen with confidence. Decision is strategic (vendor preference, OSS vs commercial, team capacity).

ABOUT THESE PATHS

Agentic Platform Engineering is the destination — an internal developer platform with self-service golden paths, AI-augmented workflows, and autonomous operations. It admits two implementation paths:

- **Three Horizons (Red Hat Developer Hub).** Three Horizons (with Red Hat Developer Hub): commercial, OpenShift-aligned, fully supported. Best fit for organizations already in the Red Hat ecosystem or requiring vendor-backed enterprise support.
- **Open Horizons (Backstage).** Open Horizons (with Backstage): open-source, vendor-neutral, fully customizable. Best fit for organizations preferring OSS, with strong platform engineering capacity, or requiring cloud-agnostic deployment.

Both paths produce the same end-state. The choice is operational (tooling, vendor relationship, support model), not capability-based. H1/H2/H3 time phasing (Months 0-6 / 6-18 / 18-36) applies to both paths.

7. Key Evidence Sources

The scoring above is supported by evidence collected during assessment sessions. The categories below list the kinds of artifacts and observations that informed the rationale.

Documents and Architectural Artifacts

- Architectural decision records reviewed across 6 critical services
- Service catalog draft and IDP target architecture proposal
- Current DevOps tooling inventory with utilization data
- Quarterly engineering survey results

Dashboards and Tool Reports

- AI coding usage dashboard (license deployment + acceptance metrics)
- DORA metrics dashboard (portfolio level)
- SLO compliance dashboards for tier-1 services
- Cost allocation dashboards by team

Interviews and Discovery Sessions

- Engineering leadership discovery sessions (CTO, VP Eng, Pillar Leads)
- Pillar working group sessions (8 capability deep-dives)
- Developer-experience round tables (15 developers across 5 teams)

Operational Metrics

- DORA metrics: deployment frequency, lead time, change failure, MTTR
- AI suggestion acceptance and adoption telemetry
- Test coverage and mutation testing rates per repository

Governance and Policy Records

- Existing engineering policies (branch protection, security scanning)
- Compliance evidence collection records
- Risk register entries from previous quarters

8. Conclusion

Cliente Exemplo S.A. has reached an overall maturity of L1.99 (L2 — Definido) across the assessed pillars. This represents a meaningful baseline from which to drive systematic improvement toward the target state.

0 capabilities currently sit at HIGH gap priority and warrant immediate attention. Pillars analyzed: P1, P2, P3. Each capability has been scored against quantifiable evidence and validated against the framework rubric.

The companion Transformation Roadmap documents (Parts 1-4) translate this current-state picture into a phased H1/H2/H3 plan with specific initiatives, success metrics, and governance structure to close every gap identified here.